

Deep Learning: A Survey

Nathaniel Beckemeyer

1 Introduction

My initial ideas were to attempt to classify the CIFAR-10 image dataset using a convolutional neural network (CNN) and autoencoding the MNIST dataset. When I searched for how to get started, both of these tasks had already been implemented—the CIFAR-10 task was actually a TensorFlow tutorial already! So instead, in order to gain intuition for the types of problems that deep learning is able to solve, I opted to try to use deep learning on a variety of tasks rather than try to do something in-depth with a CNN or some other deep neural network.

Rather than write all of my own code from scratch, I mostly used the TensorFlow code that others had already written to experiment with 5 tasks:

1. Classifying the CIFAR-10 dataset¹ with a CNN.
2. Using a recurrent neural network (RNN) to predict words in the Penn Tree Bank (PTB)² dataset.
3. Using a wide and deep network to classify adults as having income greater than or less than \$50000.³
4. Using the ‘word2vec’ method to approximate semantic distance between words.⁴
5. Using an autoencoder to encode MNIST digits.⁵

The remainder of the report will be composed of descriptions of the experiments, their results, and brief discussions of each.

2 Experiments

2.1 CNN on the CIFAR-10 dataset

This task was fairly straightforward: Image classification. Given 60000 labeled 32x32 images, the goal was to classify new images into 10 categories: Airplane, automobile, bird, cat, deer, dog, frog, horse, ship, and truck.

¹ Alex Krizhevsky: <https://www.cs.toronto.edu/~kriz/cifar.html>

² Tomas Mikolov: <http://www.fit.vutbr.cz/~imikolov/rnnlm/simple-examples.tgz>

³ UCI Machine Learning Repository: <https://archive.ics.uci.edu/ml/datasets/adult>

⁴ Matt Mahoney: <http://mattmahoney.net/dc/text8.zip>

⁵ Ships with TensorFlow

Because we've covered CNNs in class, I won't go into detail on how they work; it suffices to provide the pipeline of this network: Inputs go through two iterations of a convolutional layer, a norm response layer, and max pooling; then, through two ReLU layers to a linear softmax (for generating logits) output layer.

I followed the TensorFlow tutorial to do this part; in the code that they provide, however, the CNN training runs over the course of 1000000 steps, but I only ran it for 2000 in order to expedite termination. Despite such a few number of steps, it classified images with 65.7% accuracy (where accuracy is defined as output where the highest likelihood class is the correct class).

2.2 RNN on the PTB dataset

2.2.1 RNNs

Next, I wanted to use a recurrent neural network to predict words in a sequential reasoning task.

The recurrent neural network were what I learned is called *Long Short Term Memory Networks*, (LSTMs) which is a special kind of RNN for remembering state for long periods of time. See Figure 1 for a visualization of the network⁶, which I will explain briefly. It helps me to think of the top line as the *state line* and the bottom one as the *observation line*, which is a combination of the previous output and the current inputs. The LSTM network consists of four




Fig. 1: A visualization of an unrolled LSTM RNN.

special layers of artificial neurons—they are special because they are not merely sequential layers, but rather they interact differently.

The leftmost layer in the diagram above is the *forget gate*, and it identifies which features from the observation line should not be included in the new state. The state line is multiplied by the output from this layer to forget the unimportant features of the output.

⁶ This image, along with my understanding of LSTMs, comes from Christopher Olah's blog at <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

The layer to its right is called the *input gate*. It identifies which features from the observation line are important to store in the current state. This gate interacts with the *candidate values*, generated by the tanh layer. These two outputs are multiplied to select the changes in important features, and added to the state line to represent the current state.

The final layer is the *output gate*. It helps select the features from the observation line that are important; these features are then multiplied by the state line (scaled by tanh to be between -1 and 1) to generate the new output. (The reason for naming the sigmoidal layers the gates is to scale the values to be between 0 and 1, which is good for measuring relative importance.)

2.2.2 Task

The RNN task is to predict next words in the PTB dataset. First, the words are embedded into a dense representation. Then, the LSTM does repeated batch gradient descent on a series of words. The loss function for a series of examples follows.

$$L \leftarrow -\frac{1}{N} \sum_{i=1}^N \ln p_{target}$$

This is the natural logarithm of the average per-word perplexity, which is the result reported for this task (and, according to the tutorial I was looking at, typical of NLP tasks in general—we used perplexity of n-gram models in the advanced AI course): I ran code on this dataset for 13 epochs (or complete runthroughs of the test set), and, ultimately, it achieved a test perplexity score of 115.3.

Unfortunately, I don't have a good grasp of what that number means; reverse-engineering the formula for perplexity implies that the average logarithm probability is -4.75, or the average probability is 0.8%. That's more useful, but I don't have a good baseline for comparison.

Ultimately, I think that this task was the most fascinating of all of the ones that I did, and perhaps because I understand it least. For instance, it is unclear to me why the input gate and candidate values could not be merged into a single tanh layer that polarizes on important values and approaches 0 on unimportant values. I suspect that it has to do with the mathematics of the backpropagation. I plan to explore this area much more in the future.

2.3 Wide & Deep Learning

This task comprised of predicting whether an individual earned more or less than \$50000 based on their corresponding census data. Intriguingly, it focused on combining a 'wide' learning model with a deep learning one. See Figure 2 for a visualization.⁷ The purpose of the wide model is to memorize interactions between many sparse features, such as the combination of age, education, and

⁷ This visualization, and the task, come from the corresponding TensorFlow tutorial: https://www.tensorflow.org/tutorials/wide_and_deep.

native country. For continuous features, or features with very many values, buckets (akin to tiling in reinforcement learning in continuous domains) or hashing were performed.




Fig. 2: A wide & deep neural network example.

The end classifier is a linear classifier that combines input from all of the wide columns the last layer from the deep neural network (which had two hidden layers with 200 and 100 nodes each). I tried to play around with the size of the networks and the sparse feature columns (i.e., adding more for the combination of multiple input features), but I never managed to get the accuracy above 85%, which was only slightly (and probably not significantly) above the baseline accuracy of the tutorial.

I had never thought of using very many columns to represent interactions between different features and memorize results before—that was an incredibly neat trick. It also made me wonder about the performance of other types of bucketizations; things like Gaussian distances, rather than buckets, for instance.

2.4 word2vec

This task was simply an experiment to try to understand how the word2vec algorithm worked. The principle idea behind this algorithm is to vectorize words to prevent sparsity by assuming that words that appear close to each other are semantically similar, and guessing at different dimensions of meaning. These dimensions are called *embeddings*.

These embeddings are calculated as the hidden layer of a neural network. In this case, the input to the network is a word, and the output from the network is the context from that word; in other words, the network tries to predict the surrounding words from a given word!

I ran this model using a vocabulary of the 50000 most common (anything else was replaced with an unknown tag) words from the dataset mentioned above, following a tutorial on the TensorFlow website. The number of contextual words was 2—one on each side of the target word (the predictor). There

were 128 embeddings. The 1000 most common words were projected from the 128 dimensions the algorithm selected into two dimensions using a technique called *t-Distributed Stochastic Neighbor Embedding* (tSNE). tSNE defines a probability distribution where numbers that are close to each other in the high dimensional space have high probability, and does the same thing for numbers in the low-dimensional space. It then minimizes the Kullback-Leibler divergence of the high-dimensional distribution from the low-dimensional distribution, using gradient descent to assign the parameters of the low-dimensional one. The result is visualized in Figure 3




Fig. 3: A wide & deep neural network example.

In the future, I would like to try this on a more domain-specific dataset, and see how well it distinguishes the various concepts within that domain. I expect that it will do well, as only the relative usage of words matters.

2.5 MNIST Autoencoding

This last task was very simple. When I searched for autoencoding MNIST, the top result was code doing just that. It's a very simple neural network, and I played with it slightly to try to change the results (and modified it to visualize the encodings). There were three hidden layers: A 256-unit layer, a 144-unit layer, and then another 256-unit layer.

Autoencoding is a technique for encoding data in neural networks where the network is trained to make its output match its input exactly; the secret lies in the front half of the network where the data is encoded. By cutting the network into two, where the center layer is the output the first network and the input of the second network, you get encoding and decoding networks, respectively.

The results of this exercise are visualized in Figure 4.




Fig. 4: The MNIST autoencoding visualization. The top layer is the 28x28 original MNIST data. The middle layer is the 12x12 encoded version, and the output layer is the 28x28 reconstruction.